Loading ≠ information: the model weights items by what they reveal, not by how strongly they load

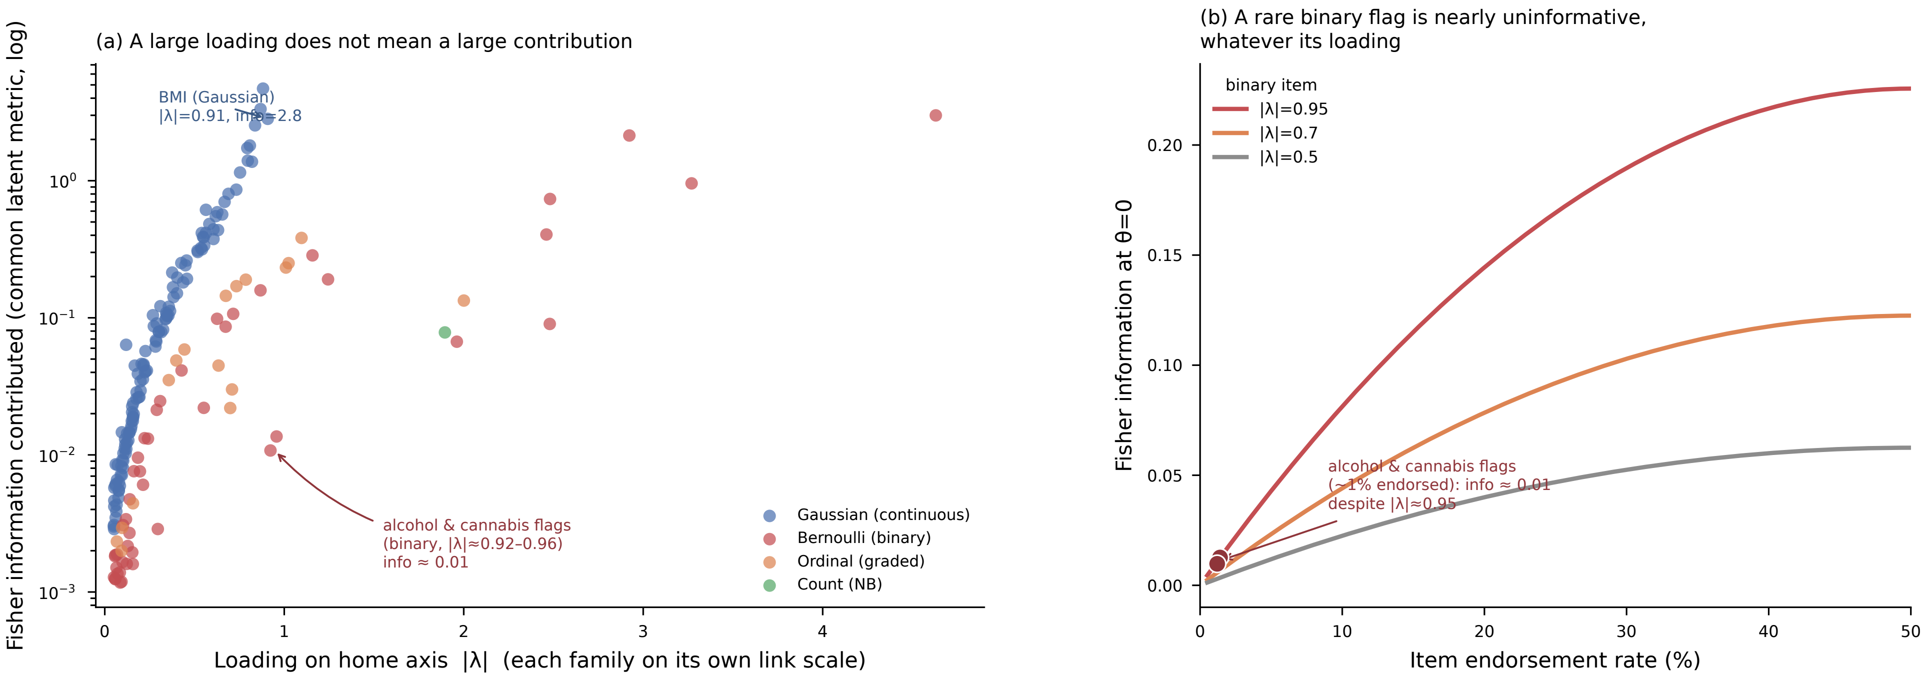

(a) Each point is one indicator: its loading on its home axis (on that likelihood family's natural link scale, so cross-family magnitudes are not directly comparable) versus the Fisher information it contributes to that axis score in the common latent metric — the quantity that IS comparable, computed exactly per family at the population mean. Continuous markers with moderate loadings dominate; two substance flags load near 0.95 yet contribute ~0.01, a hundred-fold less than a mid-loading continuous item. (b) For a binary item, information is $\lambda^2 \cdot p(1-p)$: at ~1% endorsement it collapses regardless of loading. This is why the substance axis, nominally carried by two high-loading flags, is in fact near-empty — and why an information-based reading, not a loading-based one, is the correct lens on a mixed-type measurement model.